

Optimal Stopping via Randomized Neural Networks

Calypso Herrera¹ Florian Krach¹ Pierre Ruysen²
Josef Teichmann¹

June 22th, 2021

¹Department of Mathematics, ETH Zurich

²Google Brain, Google Zurich

- 1 Introduction
- 2 Optimal Stopping via Backward Recursion
 - Least Squares Monte Carlo
 - Neural Least Squares Monte Carlo
 - Deep Optimal Stopping
 - Randomized Least Squares Monte Carlo
- 3 Reinforcement Learning
 - An alternative to the backward recursion
 - Fitted Q-learning
 - Randomized Fitted Q-learning (RFQI)
- 4 Numerical Results on Markovian Problems
 - Black Scholes - Max Call
 - Heston - Max Call
 - Geometric Put and Basket Call
 - Increasing the number of exercise dates
- 5 Non-Markovian Problems
 - Randomized Recurrent Least Squares Monte Carlo (RLSM)
- 6 Conclusion

Optimal stopping problem?

Exploration-exploitation of **standard approaches** and **machine learning based approaches** to solve the optimal stopping problem.

Least Squares Monte Carlo

- The price of the discretized American option can be expressed through the Snell envelope U .
- underlying process X , $x^i = (x_0, x_1^i, \dots, x_N^i)$ for path $i \in \{1, \dots, m\}$
- cash flow $p^i = (p_0, p_1^i, \dots, p_N^i)$ for path $i \in \{1, \dots, m\}$
- α the step-wise discounting factor
- continuation value $c_n(x) = E(\alpha U_{k+1} | X_n = x)$
- Tsitsiklis and Van Roy (2001)

$$\begin{aligned}
 p_N^i &= g(x_N^i) \\
 p_n^i &= \max \left(\underbrace{g(x_n^i)}_{\text{stop}}, \underbrace{c_n(x_n^i)}_{\text{continue}} \right)
 \end{aligned}$$

Least Squares Monte Carlo

Tsitsiklis and Van Roy 2001

$$\begin{aligned}
 p_N^i &= g(x_N^i) \\
 p_n^i &= \underbrace{g(x_n^i)}_{\text{payoff}} \underbrace{1_{\{g(x_n^i) \geq c_n(x_n^i)\}}}_{\text{exercise}} + \underbrace{c_n(x_n^i)}_{\text{continuation value}} \underbrace{1_{\{g(x_n^i) < c_n(x_n^i)\}}}_{\text{continue}}
 \end{aligned}$$

- **Longstaff and Schwartz (2001)** use the approximation only for the stopping decision.

Longstaff and Schwartz 2001

$$\begin{aligned}
 p_N^i &= g(x_N^i) \\
 p_n^i &= \underbrace{g(x_n^i)}_{\text{payoff}} \underbrace{1_{\{g(x_n^i) \geq c_n(x_n^i)\}}}_{\text{exercise}} + \underbrace{\alpha p_{n+1}^i}_{\text{discounted future price}} \underbrace{1_{\{g(x_n^i) < c_n(x_n^i)\}}}_{\text{continue}}
 \end{aligned}$$

Linear combination of features for the continuation value

Longstaff and Schwartz 2001

$$\begin{aligned}
 p_N^i &= g(x_N^i) \\
 p_n^i &= \underbrace{g(x_n^i)}_{\text{payoff}} \underbrace{1_{\{g(x_n^i) \geq c_n(x_n^i)\}}}_{\text{exercise}} + \underbrace{\alpha p_{n+1}^i}_{\text{discounted future price}} \underbrace{1_{\{g(x_n^i) < c_n(x_n^i)\}}}_{\text{continue}}
 \end{aligned}$$

- The continuation value is approximated by linear combination of basis functions,

$$c_n(x) = E(\alpha U_{n+1} | x_n = x) \approx \sum_{i=1}^n \theta_i f_i(x) = \theta^\top f(x)$$

- where the parameters θ are found by minimizing the loss function

$$\varphi(\theta_n) = \sum_{i=1}^N (c_{\theta_n}(x_n^i) - \alpha p_{n+1}^i)^2$$

- using a linear regression **at each date** $n \in \{N-1, \dots, 1\}$.

Neural networks approximation of the continuation value

Kohler et. al. (2010) = Tsitsiklis Van Roy(2001) + neural networks

$$\begin{aligned}
 p_N^i &= g(x_N^i) \\
 p_n^i &= \underbrace{g(x_n^i)}_{\text{payoff}} \underbrace{1_{\{g(x_n^i) \geq c_{\theta_n}(x_n^i)\}}}_{\text{exercise}} + \underbrace{c_{\theta_n}(x_n^i)}_{\text{continuation value}} \underbrace{1_{\{g(x_n^i) < c_{\theta_n}(x_n^i)\}}}_{\text{continue}}
 \end{aligned}$$

- The continuation value is now approximated by a neural network

$$c_k(x) \approx c_{\theta_n}(x)$$

- where the parameters θ_n are found by minimizing the loss function

$$\varphi(\theta_n) = \sum_{i=1}^N (c_{\theta_n}(x_n^i) - \alpha p_{n+1}^i)^2.$$

- using a gradient descent method at **each date** $n \in \{N-1, \dots, 1\}$.

Neural networks approximation of the continuation value

Lapeyre Lelong (2019) = Longstaff Schwartz (2001) + neural networks

$$\begin{aligned}
 p_N^i &= g(x_N^i) \\
 p_n^i &= \underbrace{g(x_n^i)}_{\text{payoff}} \underbrace{1_{\{g(x_n^i) \geq c_{\theta_n}(x_n^i)\}}}_{\text{exercise}} + \underbrace{\alpha p_{n+1}^i}_{\text{discounted future price}} \underbrace{1_{\{g(x_n^i) < c_{\theta_n}(x_n^i)\}}}_{\text{continue}}
 \end{aligned}$$

- The continuation value is now approximated by a neural network

$$c_k(x) \approx c_{\theta_n}(x)$$

- where the parameters θ_n are found by minimizing the loss function

$$\varphi(\theta_n) = \sum_{i=1}^N (c_{\theta_n}(x_n^i) - \alpha p_{n+1}^i)^2.$$

- using a gradient descent method at **each time** $n \in \{N-1, \dots, 1\}$.

Neural networks approximation of the indicator function

Becker et. al. (2019) = Longstaff Schwartz (2001) + NN for indicator

$$\begin{aligned}
 p_N^i &= g(x_N^i) \\
 p_n^i &= \underbrace{g(x_n^i)}_{\text{payoff}} \underbrace{1_{\{g(x_n^i) \geq c_k(x_n^i)\}}}_{f_{\theta_n}(x)} + \underbrace{\alpha p_{n+1}^i}_{\text{discounted future price}} \underbrace{1_{\{g(x_n^i) < c_{\theta_n}(x_n^i)\}}}_{1 - f_{\theta_n}(x)}
 \end{aligned}$$

- The indicator function is now approximated by a neural network

$$1_{\{g(x_n^i) \geq c_k(x_n^i)\}} \approx f_{\theta_n}(x)$$

- where the parameters θ_n are found by ~~minimizing~~ maximizing the loss function

$$\varphi(\theta_n) = \sum_{i=1}^N \alpha p_k^i.$$

- using a gradient descent method at **each time** $n \in \{N-1, \dots, 1\}$.

Randomized Least Squares Monte Carlo (RLSM)

Randomized Least Squares Monte Carlo (RLSM)

$$\begin{aligned}
 p_N^i &= g(x_N^i) \\
 p_n^i &= \underbrace{g(x_n^i)}_{\text{payoff}} \underbrace{\mathbf{1}_{\{g(x_n^i) \geq c_{\theta}(x_n^i)\}}}_{\text{exercise}} + \underbrace{\alpha p_{n+1}^i}_{\text{discounted future price}} \underbrace{\mathbf{1}_{\{g(x_n^i) < c_{\theta}(x_n^i)\}}}_{\text{continue}}
 \end{aligned}$$

- For each date n , going backward, the continuation value is approximated by a neural network

$$c_{\theta_n}(x) = A_n^T \sigma(Ax + b) + b_n$$

- where the parameters of the hidden layer (A, b) are randomly chosen and not optimized
- and only the parameters of the last layer $\theta_n = (A_n, b_n)$ are optimized by minimizing the loss function

$$\psi_n(\theta_n) = \sum_{i=1}^m (c_{\theta_n}(x_n^i) - \alpha p_{n+1}^i)^2. \quad (1)$$

Randomized Least Squares Monte Carlo (RLSM)

Algorithm 1 Optimal stopping via randomized least squares Monte Carlo (RLSM)

Input: discount factor α , initial value x_0

Output: price p_0

1: sample a random matrix $A \in \mathbb{R}^{(K-1) \times d}$ and a random vector $b \in \mathbb{R}^{K-1}$

2: simulate $2m$ paths of the underlying process (x_1^i, \dots, x_N^i) for $i \in \{1, \dots, 2m\}$

3: for each path $i \in \{1, \dots, 2m\}$, set $p_N^i = g(x_N^i)$

4: for each date $n \in \{N-1, \dots, 1\}$

 a: for each path $i \in \{1, \dots, 2m\}$, set $\phi(x_n^i) = (\sigma(Ax_n^i + b)^\top, 1)^\top \in \mathbb{R}^K$

 b: set $\theta_n = \alpha \left(\sum_{i=1}^m \phi(x_n^i) \phi(x_n^i)^\top \right)^{-1} \left(\sum_{i=1}^m \phi(x_n^i) p_{n+1}^i \right)$

 c: for each path $i \in \{1, \dots, 2m\}$, set $p_n^i = g(x_n^i) \mathbf{1}_{g(x_n^i) \geq \theta_n^\top \phi(x_n^i)} + \alpha p_{n+1} \mathbf{1}_{g(x_n^i) < \theta_n^\top \phi(x_n^i)}$

5: set $p_0 = \max(g(x_0), \frac{1}{m} \sum_{i=m+1}^{2m} \alpha p_1^i)$

Theorem (informal)

As the number of sampled paths m and the number of random basis functions K go to ∞ , the price p_0 computed with Algorithm 1 converges to the correct price of the Bermudan option.

An alternative to the backward recursion

- Instead of having a different function approximator for each date, **we use the same function approximator for all dates**

$$C_n(x_n) \approx C_\theta(n, x_n) \quad \forall n \in \{1, \dots, N-1\}$$

- The algorithm is initialized with a parameter vector θ_0 , and **after one sample path the parameter vector is updated** following

$$\theta_{n+1} = \theta_n + \alpha \nabla \varphi(\theta)$$

- This is the principle of **Reinforcement Learning**.
- We first make some random optimal stopping decisions and then we reinforce our learning over the iterations.
- **This is not new!** The first who introduce the reinforcement learning for the optimal stopping were **Tsitsiklis and Van Roy (1999)**.

Fitted Q-learning (FQI)

Tsitsiklis Van Roy(1999) = Off-Policy Q-learning + linear approximator

$$\begin{aligned}
 p_N^i &= g(x_N^i) \\
 p_n^i &= \underbrace{g(x_n^i)}_{\text{Payoff}} \underbrace{1_{\{g(x_n^i) \geq c_\theta(x_n^i)\}}}_{\text{exercise}} + \underbrace{c_\theta(x_n^i)}_{\text{continuation value}} \underbrace{1_{\{g(x_n^i) < c_\theta(x_n^i)\}}}_{\text{continue}}
 \end{aligned}$$

- The continuation value is approximated by a linear combination of basis functions $c_\theta(n, x_n) \approx \theta^T f(n, x)$
- Where the parameters θ are found by minimizing the loss

$$\varphi(\theta) = \sum_{i=1}^m \sum_{n=1}^N (c_\theta(n, x_n^i) - \alpha p_{n+1}^i)^2$$

- using **only one** linear regression for all dates.

Randomized Fitted Q-learning (RFQI)

$$\begin{aligned}
 p_N^i &= g(x_N^i) \\
 p_n^i &= \underbrace{g(x_n^i)}_{\text{payoff}} \underbrace{\mathbf{1}_{\{g(x_n^i) \geq c_\theta(x_n^i)\}}}_{\text{exercise}} + \underbrace{c_\theta(x_n^i)}_{\text{continuation value}} \underbrace{\mathbf{1}_{\{g(x_n^i) < c_\theta(x_n^i)\}}}_{\text{continue}}
 \end{aligned}$$

- As for reinforcement learning algorithms, we first fix the parameters θ^0 and then we update them $\theta^1, \theta^2, \dots$ through multiples iterations.
- For all dates n , the continuation value is approximated by a neural network

$$c_\theta(\tilde{x}) = A_2^\top \sigma(A\tilde{x} + b) + b_2$$

- where $\tilde{x}_n = (n, x_n)$.
- where the parameters of the hidden layer (A, b) are randomly chosen and not optimized
- and only the parameters of the last layer $\theta = (A_2, b_2)$ are optimized by minimizing the loss function

$$\psi(\theta) = \sum_{n=1}^N \sum_{i=1}^m (c_{\theta_n}(\tilde{x}_n^i) - \alpha p_{n+1}^i)^2. \quad (2)$$

Algorithm 2 Optimal stopping via randomized fitted Q-Iteration (RFQI)**Input:** discount factor α , initial value x_0 **Output:** price p_0

- 1: sample a random matrix $A \in \mathbb{R}^{(K-1) \times (d+2)}$ and a random vector $b \in \mathbb{R}^{K-1}$
- 2: simulate $2m$ paths of the underlying process (x_1^i, \dots, x_N^i) for $i \in \{1, \dots, 2m\}$
- 3: initialize weights $\theta = 0 \in \mathbb{R}^K$
- 4: for each iteration ℓ until convergence of θ
 - a: for each path $i \in \{1, \dots, 2m\}$,
 - i: set $p_N^i = g(x_N^i)$
 - ii: for each date $n \in \{1, \dots, N-1\}$,
 - set $\phi(n, x_n^i) = (\sigma(A\tilde{x}_n^i + b), 1) \in \mathbb{R}^K$
 - set $p_n^i = \max(g(x_n^i), \phi(n, x_n^i)^\top \theta)$
 - b: set $\theta = \alpha \left(\sum_{n=1}^N \sum_{i=1}^m \phi(n, x_n^i) \phi(n, x_n^i)^\top \right)^{-1} \cdot \left(\sum_{n=1}^N \sum_{i=1}^m \phi(n, x_n^i) p_{n+1}^i \right) \in \mathbb{R}^K$
- 5: set $p_0 = \max(g(x_0), \frac{1}{m} \sum_{i=m+1}^{2m} \alpha p_1^i)$

Theorem (informal)

As the number of iterations L , the number of sampled paths m and the number of random basis functions K go to ∞ , the price p_0 computed with Algorithm 2 converges to the correct price of the Bermudan option.

d	x_0	price						duration					
		LSM	DOS	NLSM	RLSM	FQI	RFQI	LSM	DOS	NLSM	RLSM	FQI	RFQI
5	80	5.80 (0.06)	5.96 (0.05)	5.86 (0.13)	5.64 (0.05)	6.08 (0.13)	6.07 (0.11)	15s	22s	9s	6s	8s	6s
	100	26.15 (0.16)	26.82 (0.20)	26.35 (0.19)	25.36 (0.16)	26.78 (0.14)	26.68 (0.25)	15s	23s	13s	6s	8s	6s
	120	50.87 (0.13)	51.50 (0.25)	51.00 (0.17)	49.89 (0.29)	51.70 (0.20)	51.67 (0.17)	15s	23s	13s	6s	8s	6s
10	80	9.88 (0.09)	10.29 (0.12)	9.93 (0.19)	9.63 (0.07)	10.40 (0.12)	10.36 (0.10)	30s	23s	9s	6s	12s	7s
	100	35.12 (0.11)	35.82 (0.14)	35.10 (0.25)	34.18 (0.15)	35.92 (0.13)	35.92 (0.11)	30s	23s	12s	6s	12s	7s
	120	61.63 (0.22)	62.64 (0.21)	61.75 (0.21)	60.67 (0.16)	62.72 (0.22)	62.77 (0.13)	30s	23s	12s	6s	14s	7s
50	80	23.03 (0.09)	24.46 (0.11)	22.72 (0.31)	22.85 (0.07)	24.59 (0.13)	24.47 (0.07)	7m59s	27s	12s	7s	5m25s	8s
	100	53.56 (0.09)	55.00 (0.12)	52.45 (0.41)	53.13 (0.11)	55.34 (0.12)	55.08 (0.14)	7m58s	26s	13s	7s	6m 9s	8s
	120	83.90 (0.20)	85.76 (0.12)	83.12 (0.72)	83.36 (0.15)	85.88 (0.13)	85.79 (0.20)	8m 1s	26s	13s	7s	6m22s	8s
100	80	26.16 (0.09)	30.47 (0.16)	27.64 (0.55)	29.13 (0.10)	30.29 (0.09)	30.58 (0.12)	35m39s	31s	13s	8s	1h17m59s	9s
	100	57.89 (0.14)	62.64 (0.25)	59.04 (0.72)	60.93 (0.10)	62.27 (0.21)	62.79 (0.19)	36m59s	31s	14s	8s	1h18m15s	9s
	120	89.00 (0.15)	94.64 (0.22)	90.20 (1.19)	92.65 (0.04)	94.35 (0.14)	94.91 (0.15)	37m 0s	31s	14s	8s	1h16m15s	9s
500	80	-	42.85 (0.14)	39.18 (1.12)	42.90 (0.05)	-	44.03 (0.08)	-	1m25s	27s	14s	-	15s
	100	-	78.10 (0.14)	73.59 (1.11)	78.22 (0.09)	-	79.50 (0.12)	-	1m41s	30s	14s	-	15s
	120	-	113.27 (0.26)	107.55 (1.26)	113.37 (0.12)	-	115.18 (0.17)	-	1m44s	31s	14s	-	15s
1000	80	-	47.91 (0.12)	44.85 (1.07)	48.56 (0.07)	-	49.66 (0.04)	-	2m58s	45s	22s	-	23s
	100	-	84.43 (0.15)	80.68 (1.20)	85.26 (0.08)	-	86.56 (0.11)	-	2m59s	42s	22s	-	23s
	120	-	121.03 (0.12)	116.82 (1.07)	121.88 (0.15)	-	123.36 (0.21)	-	2m58s	42s	22s	-	22s
2000	80	-	51.18 (0.15)	50.71 (0.62)	54.10 (0.08)	-	55.06 (0.07)	-	6m18s	1m10s	38s	-	37s
	100	-	88.53 (0.13)	87.82 (0.79)	92.19 (0.07)	-	93.29 (0.10)	-	6m17s	1m14s	39s	-	38s
	120	-	125.88 (0.21)	125.24 (0.84)	130.21 (0.12)	-	131.53 (0.07)	-	6m 7s	1m11s	38s	-	37s

Table: Max call option on Black Scholes for different number of stocks d and varying initial stock price x_0 . RFQI achieves the highest prices while being the fastest and having considerably less trainable parameters.

LSM: Least Squares Monte Carlo (Longstaff and Schwartz, 2001)

DOS: Deep Optimal Stopping (Becker et al., 2019)

NLSM: Neural Least Squares Monte Carlo (Lapeyre and Lelong, 2019)

RLSM: Randomized Least Squares Monte Carlo

FQI: Fitted Q-Iteration (Tsitsiklis and Van Roy, 2001)

RFQI: Randomized Fitted Q-Iteration

Heston - Max Call

d	price						duration					
	LSM	DOS	NLSM	RLSM	FQI	RFQI	LSM	DOS	NLSM	RLSM	FQI	RFQI
5	9.81 (0.05)	10.06 (0.09)	9.95 (0.09)	9.49 (0.08)	10.11 (0.07)	10.09 (0.12)	32s	40s	29s	24s	25s	23s
10	13.33 (0.05)	13.70 (0.11)	13.26 (0.14)	12.98 (0.08)	13.75 (0.06)	13.74 (0.07)	48s	42s	32s	25s	32s	25s
50	21.18 (0.04)	21.88 (0.06)	20.15 (0.34)	21.09 (0.03)	22.05 (0.06)	22.02 (0.06)	8m37s	59s	44s	38s	6m43s	39s
100	22.73 (0.07)	25.38 (0.11)	22.75 (0.43)	24.66 (0.07)	25.43 (0.06)	25.56 (0.06)	38m44s	1m17s	1m 1s	55s	1h18m15s	56s
500	-	33.43 (0.10)	30.52 (0.48)	33.19 (0.06)	-	34.04 (0.04)	-	4m18s	3m15s	3m 3s	-	3m 4s
1000	-	36.64 (0.06)	34.67 (0.29)	36.87 (0.05)	-	37.68 (0.07)	-	8m14s	5m55s	5m47s	-	6m 0s
2000	-	39.20 (0.08)	38.37 (0.38)	40.66 (0.06)	-	41.47 (0.06)	-	16m11s	11m15s	10m51s	-	10m41s

Table: Max call option on Heston for different number of stocks d . RFQI achieves the highest prices while being the fastest and having considerably less trainable parameters.

LSM: Least Squares Monte Carlo (Longstaff and Schwartz, 2001)

DOS: Deep Optimal Stopping (Becker et al., 2019)

NLSM: Neural Least Squares Monte Carlo (Lapeyre and Lelong, 2019)

RLSM: Randomized Least Squares Monte Carlo

FQI: Fitted Q-Iteration (Tsitsiklis and Van Roy, 2001)

RFQI: Randomized Fitted Q-Iteration

Geometric Put and Basket Call

payoff	d	price						duration					
		LSM	DOS	NLSM	RLSM	FQI	RFQI	LSM	DOS	NLSM	RLSM	FQI	RFQI
GeometricPut	1	7.08 (0.11)	6.97 (0.08)	7.05 (0.11)	7.06 (0.06)	6.92 (0.13)	6.96 (0.10)	8s	20s	9s	5s	6s	6s
	5	3.35 (0.04)	3.33 (0.03)	3.32 (0.05)	3.29 (0.05)	3.35 (0.05)	3.34 (0.03)	15s	22s	10s	6s	8s	7s
	10	2.38 (0.04)	2.37 (0.02)	2.35 (0.04)	2.35 (0.03)	2.40 (0.03)	2.40 (0.05)	29s	23s	10s	7s	13s	7s
	20	1.67 (0.01)	1.69 (0.03)	1.59 (0.05)	1.64 (0.03)	1.72 (0.02)	1.72 (0.02)	1m24s	23s	10s	7s	38s	8s
BasketCall	5	4.61 (0.08)	4.58 (0.06)	4.44 (0.07)	4.61 (0.05)	4.62 (0.04)	4.65 (0.03)	15s	22s	10s	6s	8s	7s
	10	3.58 (0.03)	3.60 (0.05)	3.38 (0.05)	3.61 (0.05)	3.61 (0.06)	3.63 (0.03)	30s	23s	10s	6s	21s	7s
	50	2.05 (0.03)	2.35 (0.02)	1.66 (0.07)	2.00 (0.02)	2.36 (0.02)	2.37 (0.02)	7m59s	27s	12s	7s	5m40s	8s
	100	1.18 (0.01)	2.11 (0.02)	1.24 (0.04)	1.78 (0.01)	2.10 (0.02)	2.12 (0.02)	40m34s	31s	14s	8s	1h17m19s	9s
	500	-	1.92 (0.01)	0.86 (0.02)	1.71 (0.01)	-	1.95 (0.01)	-	1m33s	30s	15s	-	15s
	1000	-	1.91 (0.01)	0.80 (0.02)	1.77 (0.00)	-	1.95 (0.01)	-	2m56s	48s	22s	-	22s
	2000	-	1.85 (0.00)	0.77 (0.01)	1.84 (0.00)	-	1.96 (0.00)	-	5m17s	1m22s	39s	-	37s

Table: Geometric put and basket call options on Black Scholes for different number of stocks d . RFQI achieves the highest prices while being the fastest and having considerably less trainable parameters.

LSM: Least Squares Monte Carlo (Longstaff and Schwartz, 2001)

DOS: Deep Optimal Stopping (Becker et al., 2019)

NLSM: Neural Least Squares Monte Carlo (Lapeyre and Lelong, 2019)

RLSM: Randomized Least Squares Monte Carlo

FQI: Fitted Q-Iteration (Tsitsiklis and Van Roy, 2001)

RFQI: Randomized Fitted Q-Iteration

Increasing the number of exercise dates

d	N	price					duration				
		LSM	DOS	NLSM	RLSM	RFQI	LSM	DOS	NLSM	RLSM	RFQI
10	50	34.54 (0.10)	36.06 (0.16)	35.95 (0.17)	33.13 (0.21)	35.98 (0.14)	2m45s	1m50s	1m 8s	33s	37s
	100	34.36 (0.10)	36.11 (0.14)	35.97 (0.20)	32.75 (0.25)	35.83 (0.15)	5m23s	3m55s	2m20s	1m 7s	1m13s
50	50	53.14 (0.08)	55.87 (0.14)	55.21 (0.21)	52.29 (0.08)	55.64 (0.13)	44m28s	2m 5s	1m15s	37s	44s
	100	52.77 (0.11)	56.08 (0.14)	55.60 (0.13)	51.81 (0.09)	55.66 (0.13)	1h30m27s	4m35s	2m34s	1m14s	1m29s
100	50	-	63.81 (0.14)	62.54 (0.22)	60.42 (0.12)	63.58 (0.13)	-	2m53s	1m25s	42s	48s
	100	-	64.10 (0.12)	63.22 (0.12)	59.92 (0.07)	63.58 (0.07)	-	5m57s	2m50s	1m25s	1m37s
500	50	-	80.86 (0.07)	76.45 (0.62)	78.51 (0.07)	81.11 (0.07)	-	8m50s	2m47s	1m16s	1m20s
	100	-	81.26 (0.11)	78.07 (0.42)	78.15 (0.11)	81.21 (0.10)	-	17m36s	5m37s	2m34s	2m40s

Table: Max call option on Black Scholes for different number of stocks d and higher number of exercise dates N . RFQI achieves similar prices as DOS while being the fastest and having considerably less trainable parameters.

LSM: Least Squares Monte Carlo (Longstaff and Schwartz, 2001)

DOS: Deep Optimal Stopping (Becker et al., 2019)

NLSM: Neural Least Squares Monte Carlo (Lapeyre and Lelong, 2019)

RLSM: Randomized Least Squares Monte Carlo

FQI: Fitted Q-Iteration (Tsitsiklis and Van Roy, 2001)

RFQI: Randomized Fitted Q-Iteration

Randomized Recurrent Least Squares Monte Carlo (RLSM)

$$\begin{aligned}
 p_N^i &= g(x_N^i) \\
 p_n^i &= \underbrace{g(x_n^i)}_{\text{payoff}} \underbrace{\mathbf{1}_{\{g(x_n^i) \geq c_\theta(x_n^i)\}}}_{\text{exercise}} + \underbrace{\alpha p_{n+1}^i}_{\text{discounted future price}} \underbrace{\mathbf{1}_{\{g(x_n^i) < c_\theta(x_n^i)\}}}_{\text{continue}}
 \end{aligned}$$

- The continuation value is approximated by a **recurrent neural network**

$$\begin{cases} h_n &= \sigma(A_x x_n + A_h h_{n-1} + b) \\ c_{\theta_n}(h_n) &= A_n^\top h_n + b_n \end{cases}$$

- where the parameters of the hidden state (A_x, A_h, b) are randomly chosen and not optimized
- and only the parameters of the readout map $\theta_n = (A_n, b_n)$ are optimized by minimizing the loss function

$$\psi_n(\theta_n) = \sum_{i=1}^m (c_{\theta_n}(x_n^i) - \alpha p_{n+1}^i)^2.$$

- We go forward for computing h_n and backward for computing p_n .

Fractional Brownian Motion




Figure: Payoff identity. Algorithms processing path information outperform.

DOS: Deep Optimal Stopping (Becker et al., 2019)

pathDOS: Deep Optimal Stopping using the entire path

pathDos-paper: values reported in the paper

RLSM: Randomized Least Squares Monte Carlo

RRLSM: Randomized Recurrent Least Squares Monte Carlo

Fractional Brownian Motion




Figure: Reinforcement learning does not work well in non-Markovian case.

RFQI: Randomized Fitted Q-Iteration

pathRFQI: Randomized Fitted Q-Iteration

RRFQI: Randomized Recurrent Fitted Q-Iteration

pathDos-paper: Deep Optimal Stopping using the entire path

Fractional Brownian Motion




Figure: Randomized Recurrent Least Monte Carlo (RRLSM) achieves similar results as reported in deep optimal stopping, while using only 20K paths instead of 4M for training which took only 4s instead of the reported 430s

High dimensions Hurst

payoff	d	price				duration			
		DOS	pathDOS	RLSM	RRLSM	DOS	pathDOS	RLSM	RRLSM
Identity	1	0.66 (0.01)	1.24 (0.01)	0.65 (0.02)	1.23 (0.01)	1m15s	2m34s	3s	4s
	5	1.72 (0.04)	1.59 (0.01)	2.08 (0.01)	2.15 (0.01)	1m32s	12m55s	15s	15s
	10	1.85 (0.03)	1.68 (0.01)	2.45 (0.00)	2.46 (0.00)	1m47s	26m21s	22s	24s
Max	5	0.29 (0.00)	0.52 (0.00)	0.28 (0.01)	0.52 (0.01)	1m30s	12m57s	12s	14s
	10	0.20 (0.01)	0.36 (0.00)	0.22 (0.02)	0.32 (0.02)	1m44s	25m19s	28s	21s
Mean	5	0.29 (0.00)	0.52 (0.00)	0.28 (0.01)	0.52 (0.01)	1m30s	12m57s	12s	14s
	10	0.20 (0.01)	0.36 (0.00)	0.22 (0.02)	0.32 (0.02)	1m44s	25m19s	28s	21s

Table: Identity, maximum and mean on the fractional Brownian motion with $H = 0.05$ and different number of stocks d . RRLSM achieves high prices while being much faster than pathDOS.

DOS: Deep Optimal Stopping (Becker et al., 2019)

NLSM: Neural Least Squares Monte Carlo (Lapeyre and Lelong, 2019)

RLSM: Randomized Least Squares Monte Carlo

FQI: Fitted Q-Iteration (Tsitsiklis and Van Roy, 2001)

RFQI: Randomized Fitted Q-Iteration

Conclusion

- we introduced two simple and powerful approaches:
 - **Randomized Least Squares Monte Carlo (RLSM)**
 - **Randomized Fitted Q-Iteration (RFQI)**
- They have the advantage of the state-of-the-art: they are very simple to implement and have convergence guarantees.
- They have the advantages of neural networks; they are easily scalable to high dimensions, and there is no need to choose basis functions by hand.
- On top of that, these methods are extremely fast.
- In Markovian problems, RFQI outperforms all existing methods reconfirming that reinforcement learning methods surpass backward induction methods.
- In non-Markovian problems, our **randomized recurrent neural network algorithm (RRLSM)**, achieves similar results as the path-version of DOS, while using less training data and being much faster.
- The speed of our algorithms is very promising for applications in high dimensions and with many discretization dates.

Thank you

for your attention

- Becker, S., Cheridito, P., and Jentzen, A. (2019). Deep optimal stopping. Journal of Machine Learning Research, 20:74.
- Lapeyre, B. and Lelong, J. (2019). Neural network regression for bermudan option pricing.
- Longstaff, F. A. and Schwartz, E. S. (2001). Valuing american options by simulation: a simple least-squares approach. Review of Financial Studies.
- Tsitsiklis, J. and Van Roy, B. (2001). Regression methods for pricing complex american-style options. IEEE Transactions on Neural Networks, 12(4):694–703.